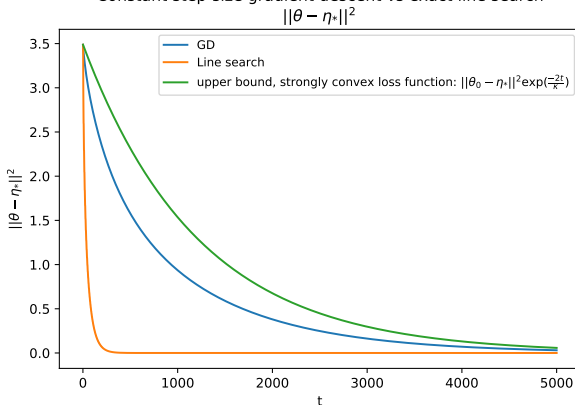


# Machine learning I, supervised learning: gradient algorithms

Constant step-size gradient descent vs exact line search



## Minimization of functions

In machine learning, we face **function minimization problems**. Typically, the function to minimize is the empirical risk on the train set, that will typically depend on a parameter  $\theta \in \mathbb{R}^d$ .

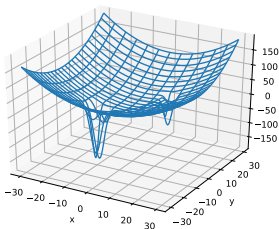


Figure – Example of a function to minimize

## Analytic minimum

What is the minimum of the function

$$f : x \rightarrow (x - 1)^2 + 3.5 \quad (1)$$

And for what value  $x$  is it obtained ?

## Context

In machine learning, we often encounter problems in high dimension, where closed-form solutions to the **empirical risk minimization** problem are **not** available (e.g. for logistic regression), or where even if they are available, the necessary computation time is too large (OLS).

## Context

In machine learning, we often encounter problems in high dimension, where closed-form solutions are not available, or where even if they are available, the necessary computation time is too large.

**Example 1** : Computing the OLS estimator requires a matrix inversion, which is  $\mathcal{O}(d^3)$ .

$$\hat{\theta} = (X^T X)^{-1} X^T y \quad (2)$$

## Context


In machine learning, we often encounter problems in high dimension, where closed-form solutions are not available, or where even if they are available, the necessary computation time is too large.

**Example 2 :** The cancellation of the gradient of the objective function with logistic loss has no closed-form solution.

## Context

Instead, we often use **iterative** algorithm such as Gradient descent (GD) or Stochastic gradient descent (SGD). SGD is the standard optimization algorithm for large-scale machine learning.

# SGD vs Lapack





## Derivation and variation

- ▶ In the case a function  $f : \mathbb{R} \rightarrow \mathbb{R}$ , we can study its variations by computing its derivative  $f'$ , **if it exists**
- ▶ If  $f'(x) > 0$ , the function grows around  $x$ .
- ▶ If  $f'(x) < 0$ , the function decreases around  $x$ .
- ▶ If  $x$  is a local extremum,  $f'(x) = 0$
- ▶ Is the reciprocal true?

## One dimensional functions

If  $f : \mathbb{R} \mapsto \mathbb{R}$  is differentiable (dérivable) :  
(Landau notation)

$$f(a + h) = f(a) + hf'(a) + o(h) \quad (3)$$

Physicist notation :

$$f(a + h) \simeq f(a) + hf'(a) \quad (4)$$

## Gradient update


In one dimension :

$$x \leftarrow x - \gamma f'(x) \quad (5)$$

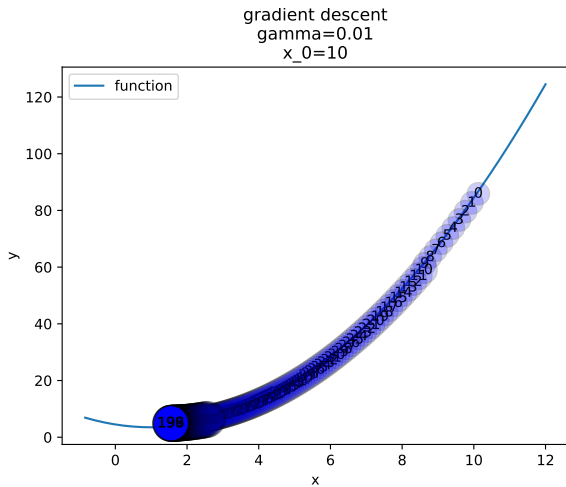
$\gamma$  is called the **learning rate** (hyperparameter of the algorithm).

## Examples in 1 dimension

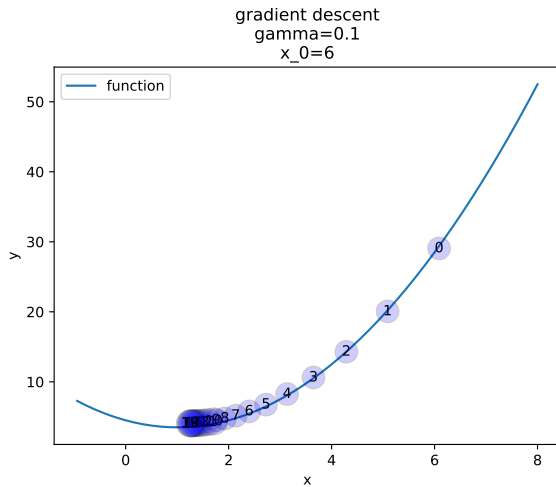
In the following examples, we use the script  
**gradients/1d\_function.py**



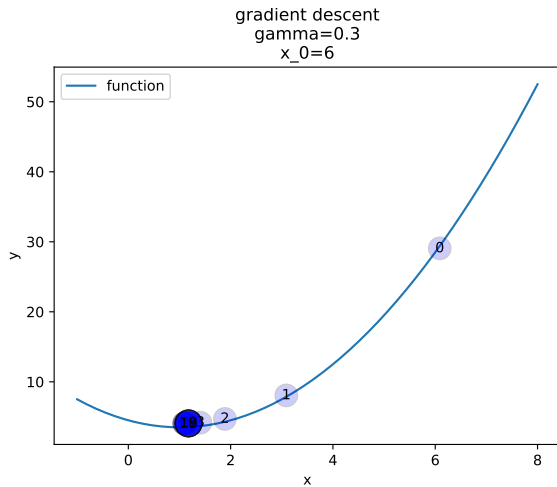
## Examples in 1 dimension



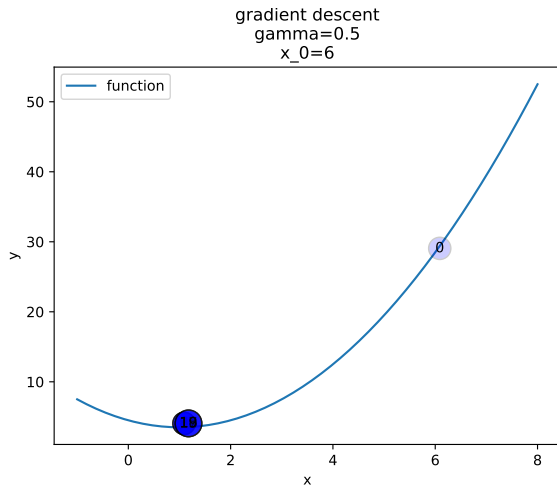
## Examples in 1 dimension



## Examples in 1 dimension




## Examples in 1 dimension






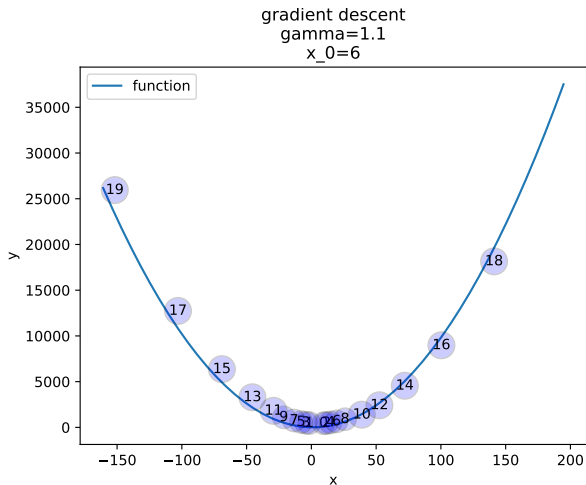
## Examples in 1 dimension



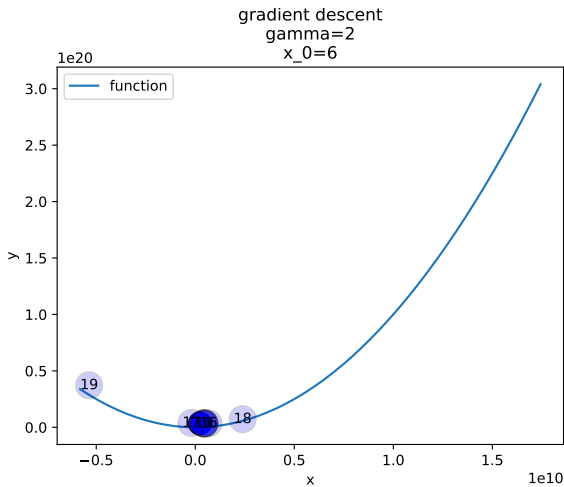
## Examples in 1 dimension



## Examples in 1 dimension



## Examples in 1 dimension



# Gradients

The **gradient** is the generalization of the derivative to functions with more than 1 input variable.

Example : consider a function  $f$  that has 2 parameters as inputs. If  $f$  is differentiable, the gradient writes :

$$\nabla f(x, y) = \left( \frac{\delta f}{\delta x}(x, y), \frac{\delta f}{\delta y}(x, y) \right) \quad (6)$$

$\frac{\delta f}{\delta x}(x, y)$  is the **partial derivative** with respect to  $x$ , computed in  $(x, y)$ .

## Example gradient

If

$$f(x, y) = x^2 + 3xy + 1 \quad (7)$$

Then

$$\forall (x, y) \in \mathbb{R}^2, \nabla f(x, y) = (2x + 3y, 3x) \quad (8)$$

## Multiple-variable functions

If  $f : \mathbb{R}^d \mapsto \mathbb{R}$  is differentiable (dérivable) :  
(Landau notation)

$$f(\theta + h) = f(\theta) + \langle \nabla f(\theta) | h \rangle + o(h) \quad (9)$$

Physicist notation :

$$f(\theta + h) \simeq f(\theta) + \langle \nabla f(\theta) | h \rangle \quad (10)$$

## Gradient descent

$$\theta \leftarrow \theta - \gamma \nabla f(\theta) \quad (11)$$

$\gamma$  must be carefully chosen.

- ▶ too large  $\gamma$  : the minimization might not work
- ▶ too small  $\gamma$  : the minimization will be too slow



## Gradient descent algorithm summary :

- ▶ In one dimension ( $x \in \mathbb{R}$ ) :

$$x \leftarrow x - \gamma f'(x) \quad (12)$$

- ▶ In  $d > 1$  dimensions,  $\theta \in \mathbb{R}^d$  :

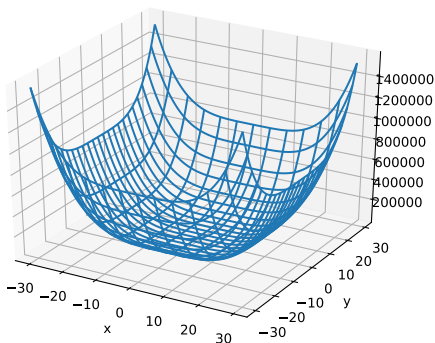
$$\theta \leftarrow \theta - \gamma \nabla f(\theta) \quad (13)$$

- ▶  $\gamma$  is the **learning rate**.

# Gradient


**Exercise 1 :** Implementing the gradient algorithm in  $\mathbb{R}^2$

We will use the algorithm on two functions defined over  $\mathbb{R}^2$ .



## Gradient

**Exercise 1 :** Implementing the gradient algorithm in  $\mathbb{R}^2$  We will use the algorithm on two functions defined over  $\mathbb{R}^2$ .

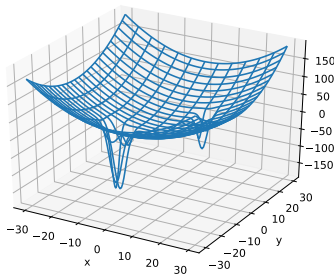



## Gradient

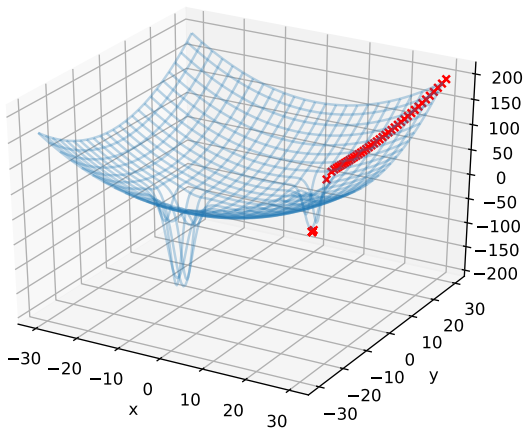
### Exercise 1 : Implementing the gradient algorithm

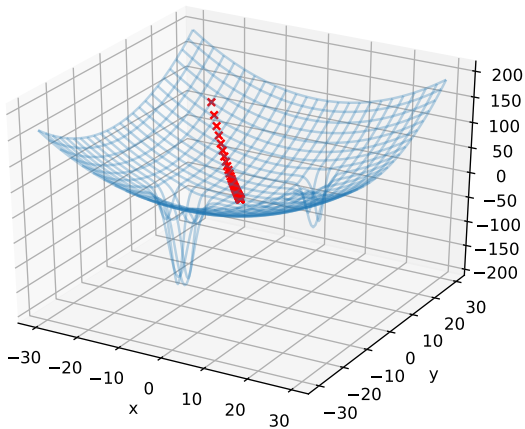
`cd ./gradient` and use the files `gradient_algo_1.py` and `gradient_algo_2.py` in order to implement the algorithm to find minima.

Experiment with all the parameters that you consider relevant (several are) to assess their impact on the algorithm.









## Convergence speed

For some problems, it is possible to have guarantees on the convergence speed of gradient descent. The results will depend on the following properties of the objective function :

- ▶ convexity or strong convexity
- ▶ smoothness (Lipshitz-continuous gradients) or non-smoothness
- ▶ condition number




## Smoothness (Example of theoretical criterion)

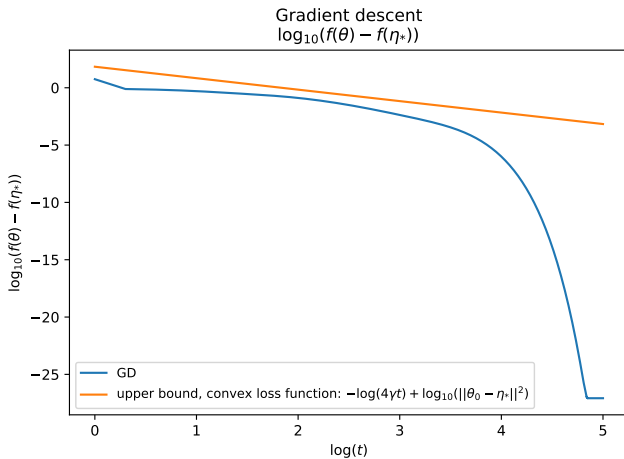
### Définition

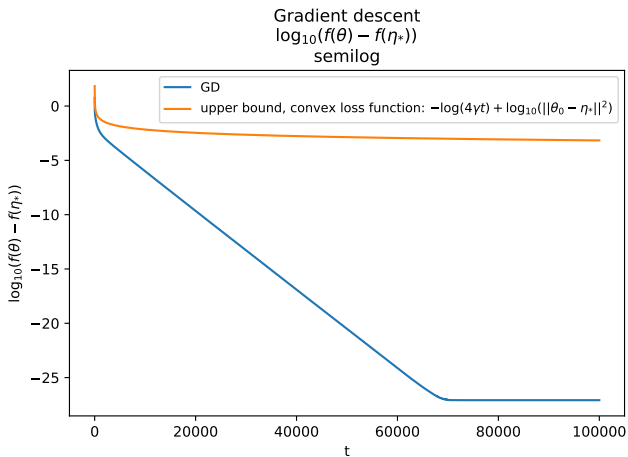
#### Smoothness

A differentiable function  $f$  with real values is said  $L$ -smooth if and only if

$$\forall x, y \in \mathbb{R}^d, |f(y) - f(x) - \nabla_x f(y - x)| \leq \frac{L}{2} \|y - x\|^2$$








## Extensions

- ▶ Line search
- ▶ Nesterov acceleration (optimal rates among algorithms that linearly combine gradients)



# Stochastic gradient descent

In machine learning, we often consider an objective function of the form

$$f(\theta) = \frac{1}{n} \sum_{i=1}^n l(y_i, f_{\theta}(x_i)) + \Omega(\theta) \quad (14)$$

## Batch gradient

In machine learning, we often consider an objective function of the form

$$f(\theta) = \frac{1}{n} \sum_{i=1}^n l(y_i, f_{\theta}(x_i)) + \Omega(\theta) \quad (15)$$

Computing the gradient of  $f$  requires at least  $n$  calculations, and each calculation also has a complexity that depends on the dimension  $d$ . When  $n$  and  $d$  are large, this can be quite slow.



## Stochastic gradient descent

We consider an objective function of the form

$$f(\theta) = \frac{1}{n} \sum_{i=1}^n l(y_i, f_{\theta}(x_i)) + \Omega(\theta) \quad (16)$$

Instead of computing the **batch gradient**  $\nabla_{\theta} f$ , we will compute :

$$u(i, \theta) = \nabla_{\theta} [l(y_i, f_{\theta}(x_i)) + \Omega(\theta)] \quad (17)$$

for a randomly sampled  $i \in [1, n]$ .

$u(i, \theta)$  is an **estimation** of the full (batch) gradient  $\nabla_{\theta} f$ .

## Tradeoff

By replacing the computation of the full gradient (GD) by this estimation (SGD) :

- ▶ we reduce the computation time (divide it by  $n$ )
- ▶ but reduce precision of the optimization

To summarize : if  $n$  is really large, SGD is often better.

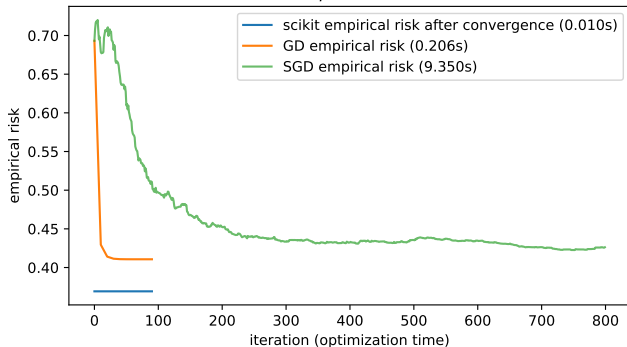
Logistic regression: loss with GD, SGD, scikit

$n=10000$ ,  $d=80$

$\gamma_{SGD0} = 0.05$ , schedule: decreasing 1

$\gamma_{GD} = 0.5$

$\mu = 0.1$



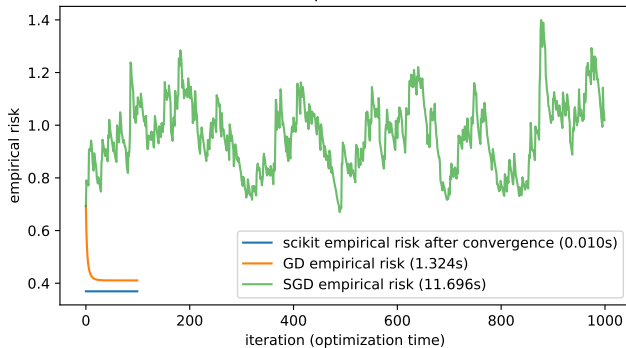
Logistic regression: loss with GD, SGD, scikit

$n=10000$ ,  $d=80$

$\gamma_{SGD0} = 0.2$ , schedule: constant

$\gamma_{GD} = 0.5$

$\mu = 0.1$



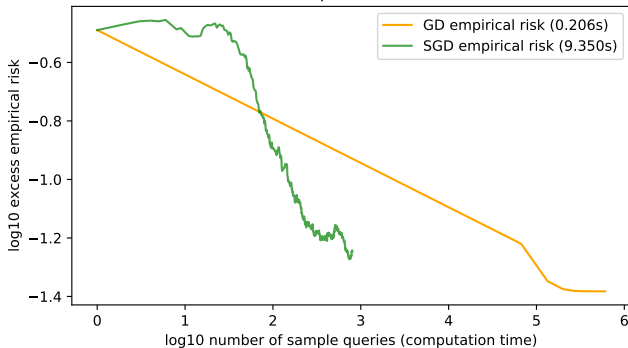
Logistic regression: log excess loss with GD, SGD, scikit

$n=10000$ ,  $d=80$

$\gamma_{SGD0} = 0.05$ , schedule: decreasing 1

$\gamma_{GD} = 0.5$

$\mu = 0.1$



## Comparison (advanced)

The ridge regression problem is smooth and strongly convex.

- ▶ GD has a convergence rate of  $\mathcal{O}(\exp(-\frac{t}{\kappa}))$ . To get an error of  $\epsilon$ , we must have  $t = \mathcal{O}(\kappa \log \frac{1}{\epsilon})$ . Since each iteration requires  $\mathcal{O}(nd)$  computations, the computation time will be  $\mathcal{O}(\kappa nd \log \frac{1}{\epsilon})$ .
- ▶ SGD has a convergence rate of  $\mathcal{O}(\frac{\kappa}{t})$ . To get an error of  $\epsilon$ , we must have  $t = \mathcal{O}(\frac{\kappa}{\epsilon})$ . Since each iteration is  $\mathcal{O}(d)$ , we have a computation time of  $\mathcal{O}(\frac{\kappa d}{\epsilon})$ .

## Comparison (advanced)

As a consequence :

- ▶ When  $n$  is large and  $\epsilon$  not too small, GD will need more computation time to reach error  $\epsilon$ . An order of magnitude can be obtained by studying the value  $\epsilon^*$  such that

$$\kappa n d \log \frac{1}{\epsilon^*} = \frac{\kappa d}{\epsilon^*}$$

Which translates to

$$\epsilon^* \log \epsilon^* = -\frac{1}{n}$$

- ▶ When  $\epsilon \rightarrow 0$ , GD becomes faster than SGD to reach this precision.

## Conclusion

For lower precision and large  $n$ , SGD is a preferable.

In machine learning, due to the estimation error that is  $\mathcal{O}(\frac{1}{\sqrt{n}})$ , a very high precision is often not needed



## Extensions of SGD

See also :

- ▶ Variance reduction methods (SAG, SAGA)